

In [12]:

```
import pandas as pd
import numpy as np
```

In []:

In [14]:

```
!pip install opencv-python
import cv2
```

Requirement already satisfied: opencv-python in c:\users\shahzad\anaconda3\new folder (2)\lib\site-packages (4.8.0.74)
Requirement already satisfied: numpy>=1.17.3 in c:\users\shahzad\anaconda3\new folder (2)\lib\site-packages (from opencv-python) (1.23.5)

In [15]:

```
from sklearn.model_selection import train_test_split
```

In [16]:

```
dataset=[]
```

FOLDER PATH

In [17]:

```
folder_path=['C:/Users/shahzad/Downloads/gender/train/female', 'C:/Users/shahzad/Downloads
```

In [18]:

```
import os
```

Data Preprocessing

In [30]:

```

# Iterate over the folder paths
for i in folder_path:
    folder_name = os.path.basename(i)

# Iterate over the images in the subdirectory
    for file_name in os.listdir(i):
        image_path = os.path.join(i, file_name)
        if os.path.isfile(image_path): # Only consider files
            # Load the image using OpenCV
            image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
            # If the image was successfully loaded
            if image is not None:
                # Resize the grayscale image to 100X100 pixels
                resized_image = cv2.resize(image, (100, 100))
                # Flatten the image and append each pixel as a separate feature along with
                flattened_image = resized_image.flatten().tolist()
                label = 'male' if 'male' in file_name else 'female' # Adjust label assign
                dataset.append(flattened_image + [label])
            else:
                print(f"Error loading image: {image_path}")

image_size = 100
num_pixels = image_size * image_size
column_names = [f'pixel_{i+1}' for i in range(num_pixels)] + ['label']
df = pd.DataFrame(dataset, columns=column_names)

print(df.head())

```

	pixel_1	pixel_2	pixel_3	pixel_4	pixel_5	pixel_6	pixel_7	pixel_8
\								
0	42	37	32	32	35	29	22	17
1	53	57	55	51	58	64	64	64
2	12	11	9	8	9	13	15	17
3	140	148	150	159	178	191	188	174
4	32	20	12	11	1	8	18	20

	pixel_9	pixel_10	...	pixel_9992	pixel_9993	pixel_9994	pixel_9995
\							
0	17	20	...	2	1	0	1
1	68	67	...	109	109	108	106
2	19	20	...	19	23	30	46
3	166	172	...	39	28	25	25
4	17	26	...	108	103	112	108

	pixel_9996	pixel_9997	pixel_9998	pixel_9999	pixel_10000	label
0	7	15	21	21	17	male
1	103	99	94	89	85	male
2	59	48	47	48	41	male
3	29	36	39	38	35	male
4	90	94	104	96	101	male

[5 rows x 10001 columns]

In []:

In [20]:

```
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
df['label_encoded'] = label_encoder.fit_transform(df['label'])
```

In [31]:

```
label = 'male' if 'male' in file_name.lower() else 'female'
```

In [32]:

```
unique_labels = df['label'].unique()
print(unique_labels)
```

```
['male' 'female']
```

DATASET SPLITTING(train,test)

In [34]:

```
# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(df.drop(columns=['label']),df['label'])

# Print the shapes of the resulting sets
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (5585, 10000)
X_test shape: (1397, 10000)
y_train shape: (5585,)
y_test shape: (1397,)
```

MACHINE LEARNING MODELS

SVM

In [35]:

```
label_counts = df['label'].value_counts()
print(label_counts)
```

```
male      3491
female    3491
Name: label, dtype: int64
```

In []:

In []:

In [25]:

```
from sklearn.metrics import confusion_matrix, classification_report
```

In [26]:

```
from sklearn.svm import SVC
```

In [27]:

```
model_svc=SVC()
```

In [37]:

```
df.rename(columns={df.iloc[:, -1].name: 'Target'}, inplace=True)
```

In [38]:

```
#get num of rows of dataset
num_rows=len(df)
#generate permuted indices
permuted_indices=np.random.permutation(num_rows)
#generate random data
random_df=df.iloc[permuted_indices]
```

In [39]:

```
X=random_df.drop('Target',axis=1)
X=X/255
X.head()
```

Out[39]:

	pixel_1	pixel_2	pixel_3	pixel_4	pixel_5	pixel_6	pixel_7	pixel_8	pixel_9
5601	0.113725	0.109804	0.113725	0.125490	0.137255	0.149020	0.121569	0.113725	0.129412
2108	0.113725	0.145098	0.168627	0.168627	0.156863	0.152941	0.180392	0.243137	0.349021
2516	0.082353	0.074510	0.070588	0.082353	0.129412	0.207843	0.301961	0.376471	0.427451
173	0.050980	0.054902	0.062745	0.066667	0.066667	0.066667	0.062745	0.062745	0.062745
33	0.011765	0.035294	0.039216	0.043137	0.082353	0.184314	0.227451	0.172549	0.113725

5 rows × 10000 columns

In [41]:

```
encoder=LabelEncoder()
y=random_df.Target
y_encoded=encoder.fit_transform(y)
y_series=pd.Series(y_encoded,name='target')
```

In [42]:

```
from sklearn.model_selection import train_test_split
np.random.seed(42)
X_train,X_test,y_train,y_test=train_test_split(X,y_encoded,test_size=0.2,random_state=42)
```

In [43]:

```
model_svc.fit(X_train,y_train);
```

In [44]:

```
y_pred_svc=model_svc.predict(X_test)
```

In [46]:

```
from sklearn.metrics import confusion_matrix,accuracy_score
```

In [47]:

```
Accuracy_svc=accuracy_score(y_pred_svc,y_test)
print('Accuracy:',Accuracy_svc)
CR=classification_report(y_pred_svc,y_test)
print('Classification Report\n',CR)
cm=confusion_matrix(y_pred_svc,y_test)
```

Accuracy: 0.28418038654259126

Classification Report

	precision	recall	f1-score	support
0	0.10	0.18	0.13	414
1	0.49	0.33	0.39	983
accuracy			0.28	1397
macro avg	0.29	0.25	0.26	1397
weighted avg	0.37	0.28	0.31	1397

In [54]:

```
cm
```

Out[54]:

```
array([[153, 520],
       [579, 145]], dtype=int64)
```

Deployment

In [69]:

```

import matplotlib.pyplot as plt
image_path = "C:/Users/shahzad/Downloads/gender_rev2/test/male/024"
user_image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)

# Resize the image to match the input size expected by the model
resized_image = cv2.resize(user_image, (100, 100))

# Flatten the image
flattened_img = resized_image.flatten()

# Normalize the flattened image data
normalized_user_image = flattened_img / 255.0

# Convert the normalized flattened image to a NumPy array and reshape it
user_input = normalized_user_image.reshape(1, -1)

# Make a prediction using the trained model
user_prediction = model_svc.predict(user_input)
image=cv2.cvtColor(resized_image,cv.COLOR_BGR2RGB)
# Decode the predicted label
predicted_class = encoder.inverse_transform(user_prediction)[0]
plt.imshow(image)
plt.title(predicted_class)

```

 -
error Traceback (most recent call last)

Cell In[69], line 6

```

3 user_image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
5 # Resize the image to match the input size expected by the model
----> 6 resized_image = cv2.resize(user_image, (100, 100))
      7 # Flatten the image
      8 flattened_img = resized_image.flatten()

```

error: OpenCV(4.8.0) D:\a\opencv-python\opencv-python\opencv\modules\imgproc\src\resize.cpp:4062: error: (-215:Assertion failed) !ssize.empty() in function 'cv::resize'

Decision Tree

In []:

In [26]:

```
from sklearn.datasets import fetch_lfw_people
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score

# Load the LFW (Labeled Faces in the Wild) dataset
lfw_people = fetch_lfw_people(min_faces_per_person=70, resize=0.4)

X = lfw_people.images
y = lfw_people.target_names[lfw_people.target]

# Flatten the images into 1D arrays
X = X.reshape(X.shape[0], -1)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a decision tree classifier
clf = DecisionTreeClassifier()

# Train the classifier
clf.fit(X_train, y_train)

# Make predictions
y_pred = clf.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.47674418604651164

logistic regression

In [49]:

```
from sklearn.linear_model import LogisticRegression
model_log=LogisticRegression()
```

In []:

In [50]:

```
model_log.fit(X_train,y_train);
```

C:\Users\shahzad\anaconda3\New folder (2)\lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

In [51]:

```
y_pred_log=model_log.predict(X_test)
```

In [52]:

```
accuracy_log=accuracy_score(y_pred_log,y_test)
print("Accuracy Score is :",accuracy_log)
C_report=classification_report(y_pred_log,y_test)
print('Classification report:',C_report)
cm=confusion_matrix(y_pred_log,y_test)
```

Accuracy Score is : 0.2133142448103078

Classification report: precision recall f1-score support

0	0.21	0.23	0.22	673
1	0.22	0.20	0.21	724
accuracy			0.21	1397
macro avg	0.21	0.21	0.21	1397
weighted avg	0.21	0.21	0.21	1397

In [55]:

```
cm
```

Out[55]:

```
array([[153, 520],
       [579, 145]], dtype=int64)
```

In []:

KNN

In [58]:

```
from sklearn.neighbors import KNeighborsClassifier
model_KNN=KNeighborsClassifier()
model_KNN.fit(X_train,y_train)
```

Out[58]:

KNeighborsClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [59]:

```
y_pred_knn=model_KNN.predict(np.array(X_test))
```

```
C:\Users\shahzad\anaconda3\New folder (2)\lib\site-packages\sklearn\base.p
y:420: UserWarning: X does not have valid feature names, but KNeighborsCla
ssifier was fitted with feature names
  warnings.warn(
```

In [60]:

```
accuracy_knn=accuracy_score(y_pred_knn,y_test)
print("Accuracy Score is :",accuracy_knn)
C_report=classification_report(y_pred_knn,y_test)
print('Classification report:',C_report)
cm=confusion_matrix(y_pred_knn,y_test)
```

Accuracy Score is : 0.2297780959198282

```
Classification report:          precision    recall  f1-score   suppo
rt
```

0	0.21	0.23	0.22	646
1	0.26	0.23	0.24	751
accuracy			0.23	1397
macro avg	0.23	0.23	0.23	1397
weighted avg	0.23	0.23	0.23	1397

In [62]:

cm

Out[62]:

```
array([[151, 495],
       [581, 170]], dtype=int64)
```

In [65]:

```
import matplotlib.pyplot as plt

algorithms = ['SVM', 'KNN', 'Decision Tree', 'Logistic Regression']
accuracy = [Accuracy_svc, accuracy_knn, accuracy_log]

colors = ['red', 'pink', 'blue', 'black']
bar_width = 0.5 # Width of the bars
bar_positions = range(len(algorithms))

plt.bar(bar_positions, accuracy, color=colors, width=bar_width)
plt.xlabel('Algorithms')
plt.ylabel('Accuracy')
plt.title('Accuracy Comparison of Machine Learning Algorithms')
plt.ylim(0, 1) # Set y-axis limits between 0 and 1

# Annotate each bar with its accuracy value
for i, acc in enumerate(accuracy):
    plt.text(i, acc + 0.02, f'{acc:.2f}', ha='center')

# Adjust x-axis labels and positions
plt.xticks(bar_positions, algorithms, rotation=15, ha='right')

plt.tight_layout() # To prevent labels from being cut off
plt.show()
```

 -
ValueError Traceback (most recent call last)

Cell In[65], line 10

```

    7 bar_width = 0.5 # Width of the bars
    8 bar_positions = range(len(algorithms))
--> 10 plt.bar(bar_positions, accuracy, color=colors, width=bar_width)
    11 plt.xlabel('Algorithms')
    12 plt.ylabel('Accuracy')
```

File ~\anaconda3\New folder (2)\lib\site-packages\matplotlib\pyplot.py:241

2, in bar(x, height, width, bottom, align, data, **kwargs)

```

    2408 @_copy_docstring_and_deprecators(Axes.bar)
    2409 def bar(
    2410     x, height, width=0.8, bottom=None, *, align='center',
    2411     data=None, **kwargs):
-> 2412     return gca().bar(
    2413         x, height, width=width, bottom=bottom, align=align,
    2414         **({"data": data} if data is not None else {}), **kwargs)
```

File ~\anaconda3\New folder (2)\lib\site-packages\matplotlib__init__.py:1

442, in _preprocess_data.<locals>.inner(ax, data, *args, **kwargs)

```

    1439 @functools.wraps(func)
    1440 def inner(ax, *args, data=None, **kwargs):
    1441     if data is None:
-> 1442         return func(ax, *map(sanitize_sequence, args), **kwargs)
    1443     bound = new_sig.bind(ax, *args, **kwargs)
    1444     auto_label = (bound.arguments.get(label_namer)
    1445                  or bound.kwargs.get(label_namer))
```

File ~\anaconda3\New folder (2)\lib\site-packages\matplotlib\axes_axes.p

y:2417, in Axes.bar(self, x, height, width, bottom, align, **kwargs)

```

    2414     if yerr is not None:
    2415         yerr = self._convert_dx(yerr, y0, y, self.convert_yunits)
-> 2417 x, height, width, y, linewidth, hatch = np.broadcast_arrays(
    2418     # Make args iterable too.
    2419     np.atleast_1d(x), height, width, y, linewidth, hatch)
    2421 # Now that units have been converted, set the tick locations.
    2422 if orientation == 'vertical':
```

File <__array_function__ internals>:180, in broadcast_arrays(*args, **kwargs)

File ~\anaconda3\New folder (2)\lib\site-packages\numpy\lib\stride_tricks.py:540, in broadcast_arrays(subok, *args)

```

    533 # nditer is not used here to avoid the limit of 32 arrays.
    534 # Otherwise, something like the following one-liner would suffice:
    535 # return np.nditer(args, flags=['multi_index', 'zerosize_ok'],
    536 #                  order='C').itviews
    538 args = [np.array(_m, copy=False, subok=subok) for _m in args]
--> 540 shape = _broadcast_shape(*args)
    542 if all(array.shape == shape for array in args):
    543     # Common case where nothing needs to be broadcasted.
    544     return args
```


File ~\anaconda3\New folder (2)\lib\site-packages\numpy\lib\stride_tricks.py:422, in _broadcast_shape(*args)

```

    417 """Returns the shape of the arrays that would result from broadcast
    418 the
    419 supplied arrays against each other.
```

```
419 """
420 # use the old-iterator because np.nditer does not handle size 0 ar
rays
421 # consistently
--> 422 b = np.broadcast(*args[:32])
423 # unfortunately, it cannot handle 32 or more arguments directly
424 for pos in range(32, len(args), 31):
425     # ironically, np.broadcast does not properly handle np.broadca
st
426     # objects (it treats them as scalars)
427     # use broadcasting to avoid allocating the full array
```

ValueError: shape mismatch: objects cannot be broadcast to a single shape. Mismatch is between arg 0 with shape (4,) and arg 1 with shape (3,).



In []:

In []: